

Sync → prevent data race

Parallel Applications

Laboratory 4
Deadline: 2 weeks

This assignment will help us learn and evaluate

- (i) two ways of writing parallel applications – using threads and using processes,
- (ii) two ways of doing inter-process communication – using shared memory and using pipes,
- (iii) two ways of doing process synchronization – using atomic operations and using semaphores.

2 threads!
 accessing same memory for so order of their execution changes it
 lack of atomicity
 3 parallel units from lab 1 → S1, S2, S3, write

Part I: Baseline application without any parallelism

This part is easy. Use your submission from laboratory 1. Update the makefile as indicated in the instructions below.

Part II

Now suppose you have a processor with three cores. You want to use all of them to make your application finish faster. You can do this by having the first core do S1. As pixels get ready, they are passed to the second core (don't wait for the S1 of the whole image to be completed before communicating to S2). The second core does S2 and passes its results to the third core. The third core does S3 and the file writing. Do this in the following ways:

1. S1, S2, and S3 are performed by 3 different processes that communicate via pipes (or fifos) (further reading on pipes)
2. S1, S2, and S3 are performed by 3 different processes that communicate via shared memory. Synchronization is done using atomic operations.
3. S1, S2, and S3 are performed by 3 different threads of the same process. They communicate through the process' address space itself. Synchronization is done using semaphores.

- Devise a method to prove in each parallel case that the pixels were received as sent, in the sent order. Describe the method in your report.
- Study the run-time and speed-up of each of the approaches and discuss.

- It is likely that the file reading and writing times dominate, and so the speed-up obtained by using three cores is negligible. So we will modify our experiment to make it compute-intensive, instead of IO-intensive. Read the image only once, but perform the transformation 1000 times. That is, S1, S2, and S3 are done 1000 times each. The results of the first run of S1 are used by the first run of S2, and the results of the first run of S2 are used by the first run of S3. For the

pipeline
older
info of
IF going
on ~ S1, S2, S3

Should be
some speedup

pass

from S1 to S2

Send hash and check

run S1 thousand times then S2 ...
but write to file only once. to S3 change size speedup

test & set locking

key is already shared no need to share again

both checksum w/ data so receiver

hashed data (1000x) and does check w/ checksum

second run of S1, it uses the same input image that has already been read into memory. The results of the second run of S1 are used by the second run of S2, and so on. The results of the 1000th run of S3 are written to the file.

- Discuss the relative ease/ difficulty of implementing/ debugging each approach.

Submit a single zip file with the source code (organized into multiple folders, one for each question), a makefile, an input *ppm* image, and a report.

- *make part1* should compile the Part I version of the code and run it, creating the file *output_part1.ppm*
- *make part2_1* should compile the multi-process, pipe version of the code and run it, creating the file *output_part2_1.ppm*
- *make part2_2* should compile the multi-process, atomic operation version of the code and run it, creating the file *output_part2_2.ppm*
- *make part2_3* should compile the multi-thread, semaphore version of the code and run it, creating the file *output_part2_3.ppm*

for